

KnowledgeDesk

Semantic internal search for mid-size companies. Ask questions. Get answers. Stop losing institutional knowledge.

Logits.in · Product Definition Document · 2025–26

THE PROBLEM

Why Companies Bleed Time on Search

Every company between 200 and 2,000 employees accumulates years of institutional knowledge in Google Drive, Confluence, Notion, email threads, and shared folders. New hires spend 3–6 weeks just *finding* things. Senior employees re-answer the same questions weekly. The COO's process SOP lives in a doc nobody has opened in 14 months.

Research by McKinsey found employees spend 1.8 hours per day — 9.3 hours per week — searching for information. That is 23% of a working year, per person.

23% of work time lost to search	6 wks avg. new hire onboarding lag	0% keyword search understands intent
---	--	--

THE SOLUTION

What KnowledgeDesk Does

KnowledgeDesk is a semantic search layer deployed inside a company's environment. It indexes all internal documents, and lets any employee ask plain-English questions and receive cited, accurate answers — not just a list of links.

Semantic Q&A;

Employees type questions like 'What is the refund policy for enterprise clients?' and get a direct answer with the source document cited.

Universal Connector

Indexes Google Drive, Confluence, Notion, SharePoint, and local file shares. One search box for everything.

Role-Based Access

Respects existing permissions — employees only see answers from documents they have access to.

One-Click Deploy

Docker-based setup on the company's own cloud (Azure/AWS/GCP). No data leaves their environment.

Admin Dashboard

See what people are searching, find knowledge gaps, track adoption metrics.

TECHNOLOGY

Stack & Build Advantage

Search Engine	Vespa (Monorail-engine) + Qdrant vector DB — already built
Embedding	OpenAI text-embedding-3-small or local SBERT model
Backend	FastAPI (Python) — Monorail-drive architecture reused
Connectors	Google Drive API, Confluence REST, Notion API, SharePoint
Auth	SSO via OAuth2 / Azure AD — respects existing access controls
Infra	Docker Compose, one-click deploy via existing Terraform scripts
Build time est.	8–12 weeks for MVP with 2–3 connectors

Why this is 60–70% already built: The Vespa schema, vector ingestion pipeline, FastAPI wrapper, and Terraform Azure infra from Monorail-engine are directly reusable. The new work is connector integrations and the frontend UI.

How to Sell It Without Traction

This product does not need a user base. The buyer is a single company that has a painful, visible problem. The sales motion is:

- 1. Identify mid-size companies (300–1500 employees) in IT services, pharma, logistics, or finance.
- 2. Contact the CTO, VP Engineering, or Head of HR directly via LinkedIn.
- 3. Offer a live demo — ask for a sample of 50 of their own documents, run the demo on their data in 48 hours.
- 4. First contract: ■8–15L/year. No per-seat complexity. Flat annual license.
- 5. First 3 customers become case studies. Referrals follow.

REVENUE MODEL

Pricing & Path to 1 Crore

Tier 1 — Starter	Up to 500 employees, 3 connectors — ■8L/year
Tier 2 — Growth	Up to 1,500 employees, unlimited connectors — ■15L/year
Tier 3 — Enterprise	Custom deployment, SLA, dedicated support — ■25–40L/year
Setup Fee	One-time ■1–2L for deployment and onboarding
Path to 1 Cr	7 Tier-1 contracts OR 4 Tier-2 contracts OR 3 mid Enterprise deals

7 Tier-1 clients = ■56L ARR	4 Tier-2 clients = ■60L ARR	3 Enterprise = ■90L+ ARR
---------------------------------------	---------------------------------------	------------------------------------

IDEAL CUSTOMER PROFILE

Who to Target First

Company size	300–2,000 employees
Industry	IT services, pharma, BFSI, logistics, manufacturing
Geography	India-first: Ahmedabad, Pune, Bengaluru, Hyderabad
Buyer persona	CTO / VP Eng for tech fit; CHRO for business case

Pain trigger	Recent rapid hiring, remote teams, compliance audits, high onboarding cost
Disqualifier	Under 200 employees (insufficient docs) or pure product startups with flat docs